

TransitionsInTimeseries.jl: A performant, extensible and reliable software for reproducible detection and prediction of transitions in timeseries

Jan Swierczek-Jereczek¹ and George Datseris²

¹ Department of Earth Physics and Astrophysics, Complutense University of Madrid. ² Department of Mathematics and Statistics, University of Exeter.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Oskar Lavery](#)

Reviewers:

- [@felixcremer](#)
- [@sgeorge91](#)

Submitted: 23 February 2024

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Transitions of nonlinear dynamical systems can significantly impact individuals and society. Examples of this are ubiquitous and include the onset of cardiac arrhythmia (Tse, 2016), the deglaciation of Earth about 20,000 years ago (Wolff et al., 2010) and the recent price collapse of many cryptocurrencies (Ismail et al., 2020). The systems displaying such transitions are usually monitored by measuring state variables that are believed to be representative of the underlying process. Researchers analyze the resulting timeseries with various methods to detect past transitions and predict future ones.

Statement of need

Over the last decades, methods to detect and predict transitions from timeseries have gained a lot of attention, both inside and outside of the scientific community. For instance, recent work predicting a collapse of the Atlantic Meridional Overturning Circulation between 2025 and 2095 has led to no less than 870 news outlets and 4100 tweets (Ditlevsen & Ditlevsen, 2023), largely because of the substantial implications of such a collapse for human societies. A common concern in the scientific community is that published work on the topic is difficult to reproduce, despite the impact it implies for humanity. This can be largely addressed by a unifying software that is accessible, performant, reproducible, reliable and extensible. Such a software does not exist yet, but here we propose TransitionsInTimeseries.jl to fill this gap. We believe this is a major step towards establishing a software as standard, widely used by academics working on transitions in timeseries.

TransitionsInTimeseries.jl

Accessibility

Open-source

TransitionsInTimeseries.jl is a free and open-source software, written in Julia and developed on GitHub, which allows any user to track the full history of the changes made to the software as well as to suggest new ones by opening a pull request or an issue.

Ease of use

TransitionsInTimeseries.jl is accessible to any scientist thanks to the convenience functions it provides to detect and predict transitions in timeseries with only a few lines of code. A frequent

36 prediction technique relies on observing, prior to a transition, an increase of the variance and
37 the AR1 regression coefficient of the detrended timeseries, which is a consequence of Critical
38 Slowing Down (CSD, (Scheffer et al., 2009)) and is here measured by Kendall's τ coefficient.
39 To assess whether this increase is significant, one can perform a statistical test, for instance by
40 performing the same computations on 1,000 surrogates of the original timeseries (Haaga &
41 Datseris, 2022). The increase in variance and AR1 coefficient can be considered significant if
42 the original timeseries classifies in the uppermost 5% of the surrogates, corresponding to a
43 p-value $p < 0.05$. All these steps can be performed, along with a visualisation of the results
44 within a few lines only:

```
# Loading and preprocessing the data needs to be done by the user
time, data = load_data()

# Choose the indicators and how to measure their change over time
indicators = (var, ar1_whitenoise)
change_metrics = (kendalltau, kendalltau)

# Configuration with adequate parameters of the sliding window over a segment
config = SegmentedWindowConfig(indicators, change_metrics, [time[1]], [time[end]]);
width_ind = length(residual) ÷ 2, whichtime = last, min_width_cha = 100)

# Compute the metrics over sliding windows and their significance
results = estimate_changes(config, data, time)
signif = SurrogatesSignificance(n = 1000, tail = :right, rng = Xoshiro(1995))
flags = significant_transitions(results, signif)

# Visualize the results
fig = plot_changes_significance(results, signif)
```

45 We apply this code to data generated by a Ricker model presenting an abrupt transition at
46 $t = 860$, which is used in the first tutorial of ewstools (Bury, 2023), the most recent software
47 covering similar functionalities. The results are shown in Fig. 1 and display, as expected from
48 CSD theory, an increase in both variance and AR1 coefficient, which is exactly the same as
49 computed by ewstools. However, calling signif.pvalues shows that the increase in variance
50 is not significant ($p = 0.284$), whereas the increase in AR1 coefficient is ($p = 0.001$).

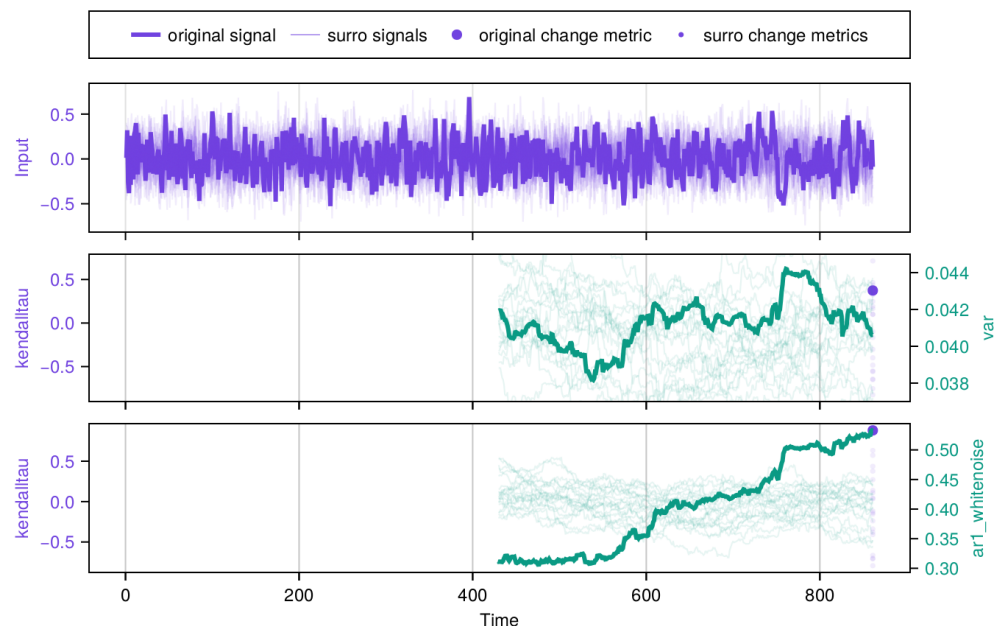


Figure 1: Output of plotting function in usage example.

We believe that a concise and unambiguous code will greatly reduce the programming effort of many researchers and ease the code reviewing process. Finally, the code documentation provides a thorough API description as well as additional examples, showcasing that the simplicity of the code also applies to real-world applications.

Performance

TransitionsInTimeseries.jl is written in Julia, which offers both a simple syntax and good performance. Additionally, all performance-relevant steps have been optimized and parallelized when possible, as, for instance, the significance testing relying on surrogates. In the final section of this article, we present a comparison to ewstools, showing that TransitionsInTimeseries.jl offers a significant speed-up in all the studied cases.

Reproducibility

Some steps of a transition analysis involve random number generators, which need to be handled with care in parallelized codes. This is done in TransitionsInTimeseries.jl, which offers the possibility of seeding a random number generator by using the keyword argument `rng`, as done in the example shown above. Furthermore, TransitionsInTimeseries.jl follows the guidelines of semantic versioning which, along with Julia's integrated package manager, ensures that the same code is used for both results generation and peer reviewing.

Reliability

In the high-impact context mentioned in the introduction, it is crucial to avoid errors. TransitionsInTimeseries.jl is therefore tested via continuous integration on a large test suite, thus providing a reliable research framework. Furthermore, a centralized code base implies that any new user is a new test, thus increasing the reliability of the code over time. Finally, the robustness of the results with respect to a parameter, e.g. the width of the sliding window, can be easily studied thanks to the simple syntax, thus contributing to the reliability of the results.

75 Extensibility

76 An important aspect of the modularity mentioned above, is that self-written functions can be
77 passed as indicators or change metrics. Thus, researchers can easily test new methods without
78 any programming overhead, nor modification of the source code. To illustrate this, the code
79 shown above can include the skewness as indicator of the transition by modifying a few lines:

```
skewness(x::Vector) = mean( (x .- mean(x))^3 ) / mean( (x .- mean(x))^2 )^1.5
indicators = (var, ar1_whitenoise, skewness)
change_metrics = (kendalltau, kendalltau, kendalltau)
```

80 There is no complexity restriction on the self-programmed functions, as long as they comply
81 with the structure of taking a vector as input and returning a scalar as output.

82 Integration

83 TransitionsInTimeseries.jl is designed to be well integrated into the Julia ecosystem. Functions
84 can be imported from other packages and subsequently passed as indicators or change metrics.
85 For instance, the skewness implemented above can be loaded from StatsBase.jl instead.
86 TransitionsInTimeseries.jl therefore offers an extremely wide and potentially unlimited library
87 of indicators. Furthermore, TimeseriesSurrogates.jl (Haaga & Datseris, 2022) is used to create
88 surrogates of the timeseries, thus offering optimized routines with numerous surrogate types.

89 Versatility

90 Choosing a pipeline

91 TransitionsInTimeseries.jl covers methods for prediction as well as detection of transitions,
92 which is unprecedented to our knowledge. This relies on the definition of different analysis
93 pipelines, which consist in a ChangesConfig determining the behavior of estimate_changes
94 via multiple dispatch. For instance, a detection task can be performed by replacing the
95 SegmentedWindowConfig by a SlidingWindowConfig in the code above:

```
# Here the data should not be detrended
time, data = load_data()

indicators = (nothing, nothing)
change_metrics = (difference_of_mean(), difference_of_max())
config = SlidingWindowConfig(indicators, change_metrics;
    width_cha = 50, whichtime = midpoint)
results = estimate_changes(config, data, time)
```

96 We here skip the computation of indicators and compare the difference in mean and maximum
97 values between the two halves of the sliding window, which gives a particularly high value
98 in the case of an abrupt transition and is therefore suited for some detection tasks. Most
99 importantly, this examples shows that the user can choose the type of analysis pipeline, along
100 with its underlying parameters (e.g. sliding window width). Similarly, different ways of testing
101 for significance are provided and can be interchangeably used.

102 Creating your own pipeline

103 Besides choosing among the already provided analysis pipelines, the user can implement
104 their own one by defining a new ChangesConfig and the corresponding behavior of
105 estimate_changes. This makes it particularly easy to leverage pre-existing functionalities of
106 TransitionsInTimeseries.jl with a minimal restriction on the structure. As explained in the
107 devdocs, the latter eases the integration of new methods into a unified framework. This also
108 holds for the significance pipeline and makes TransitionsInTimeseries.jl particularly versatile.

Comparison to already existing alternatives

earlywarnings (Dakos et al., 2012) and spatialwarnings are toolboxes written in R providing many tools to predict transitions. These are early and valuable efforts but are (1) restricted to prediction tasks, (2) written in a less performant language, (3) not parallelised, (4) not designed for convenient reproducibility and (5) not extensible.

ewstools (Bury, 2023) is a Python/TensorFlow package offering similar functionalities as earlywarnings, as well as a deep-learning approach to predicting transitions (Bury et al., 2021). This effort addresses some drawbacks of earlywarnings, mainly (4) and to a lesser extent (2) and (3). This is a great step forward but does not provide a generic and extensible framework for researchers to test new methods on both detection and prediction tasks. We believe that this is now covered by TransitionsInTimeseries.jl.

Using TransitionsInTimeseries.jl, we reproduced the computations showcased in Tutorial 1 and Tutorial 2 of the ewstools documentation, along with the block bootstrapping. We performed each computation 100 times and show the resulting run times in Fig. 2. It appears that all computation are faster in TransitionsInTimeseries, with a speed-up factor ranging from 0 to 3 orders of magnitude. The implementation of the deep-learning classifiers for transition prediction developed in (Bury et al., 2021), as well as dealing with multidimensional timeseries, are part of future developments of TransitionsInTimeseries.jl.




Figure 2: Performance comparison between ewstools and TransitionsInTimeseries.jl.

Documentation

The documentation of TransitionsInTimeseries.jl is available at <https://docs.juliahub.com/General/TransitionsInTimeseries/stable/>

Acknowledgements

Jan Swierczek-Jereczek is funded by CriticalEarth, grant no. 956170, an H2020 Research Infrastructure of the European Commission. George Datseris is funded by UKRI's Engineering and Physical Sciences Research Council, grant no. EP/Y01653X/1 (grant agreement for a EU Marie Skłodowska-Curie Postdoctoral Fellowship).

References

- Bury, T. M. (2023). Ewstools: A Python package for early warning signals of bifurcations in time series data. *Journal of Open Source Software*, 8(82), 5038. <https://doi.org/10.21105/joss.05038>
- Bury, T. M., Sujith, R. I., Pavithran, I., Scheffer, M., Lenton, T. M., Anand, M., & Bauch, C. T. (2021). Deep learning for early warning signals of tipping points. *Proceedings of the National Academy of Sciences*, 118(39), e2106140118. <https://doi.org/10.1073/pnas.2106140118>
- Dakos, V., Carpenter, S. R., Brock, W. A., Ellison, A. M., Guttal, V., Ives, A. R., Kéfi, S., Livina, V., Seekell, D. A., Nes, E. H. van, & Scheffer, M. (2012). Methods for Detecting Early Warnings of Critical Transitions in Time Series Illustrated Using Simulated Ecological Data. *PLoS ONE*, 7(7), e41010. <https://doi.org/10.1371/journal.pone.0041010>
- Ditlevsen, P., & Ditlevsen, S. (2023). Warning of a forthcoming collapse of the Atlantic meridional overturning circulation. *Nature Communications*, 14(1), 4254. <https://doi.org/10.1038/s41467-023-39810-w>
- Haaga, K. A., & Datseris, G. (2022). TimeseriesSurrogates.jl: A Julia package for generating-surrogate data. *Journal of Open Source Software*, 7(77), 4414. <https://doi.org/10.21105/joss.04414>
- Ismail, M. S., Hussain, S. I., & Noorani, M. S. Md. (2020). Detecting Early Warning Signals of Major Financial Crashes in Bitcoin Using Persistent Homology. *IEEE Access*, 8, 202042–202057. <https://doi.org/10.1109/ACCESS.2020.3036370>
- Scheffer, M., Bascompte, J., Brock, W. A., Brovkin, V., Carpenter, S. R., Dakos, V., Held, H., Nes, E. H. van, Rietkerk, M., & Sugihara, G. (2009). Early-warning signals for critical transitions. *Nature*, 461(7260), 53–59. <https://doi.org/10.1038/nature08227>
- Tse, G. (2016). Mechanisms of cardiac arrhythmias. *Journal of Arrhythmia*, 32(2), 75–81. <https://doi.org/10.1016/j.joa.2015.11.003>
- Wolff, E. W., Barbante, C., Becagli, S., Bigler, M., Boutron, C. F., Castellano, E., Angelis, M. de, Federer, U., Fischer, H., Fundel, F., Hansson, M., Hutterli, M., Jonsell, U., Karlin, T., Kaufmann, P., Lambert, F., Littot, G. C., Mulvaney, R., Röthlisberger, R., ... Wegner, A. (2010). Changes in environment over the last 800,000 years from chemical analysis of the EPICA Dome C ice core. *Quaternary Science Reviews*, 29(1–2), 285–295. <https://doi.org/10.1016/j.quascirev.2009.06.013>